Assignment 2

Assignment Date	21.09.2022
Student Name	V.Srividhya
Student Roll Number	2019115106
Maximum Marks	2 Marks

Question-1:

Download Dataset

Question-2:

Loading the dataset

Solution:

```
import pandas as pd
df=pd.read_csv('Churn_Modelling.csv')
```

Screenshot:

Question-3.1:

Univariate Analysis

Solution:

```
df.dtypes
df.head()
df.shape
import matplotlib.pyplot as plt
import seaborn as sns
plt.scatter(df.index,df['Geography'])
plt.scatter(df.index,df['NumOfProducts'])
plt.scatter(df.index,df['Age'])
import numpy as np
print('Number of users who have a credit card ',np.sum(df['HasCrCard']==1))
print('Number of users who dont have a credit card ',np.sum(df['HasCrCard']==0))
plt.hist(df['HasCrCard'])
```

Screenshot:

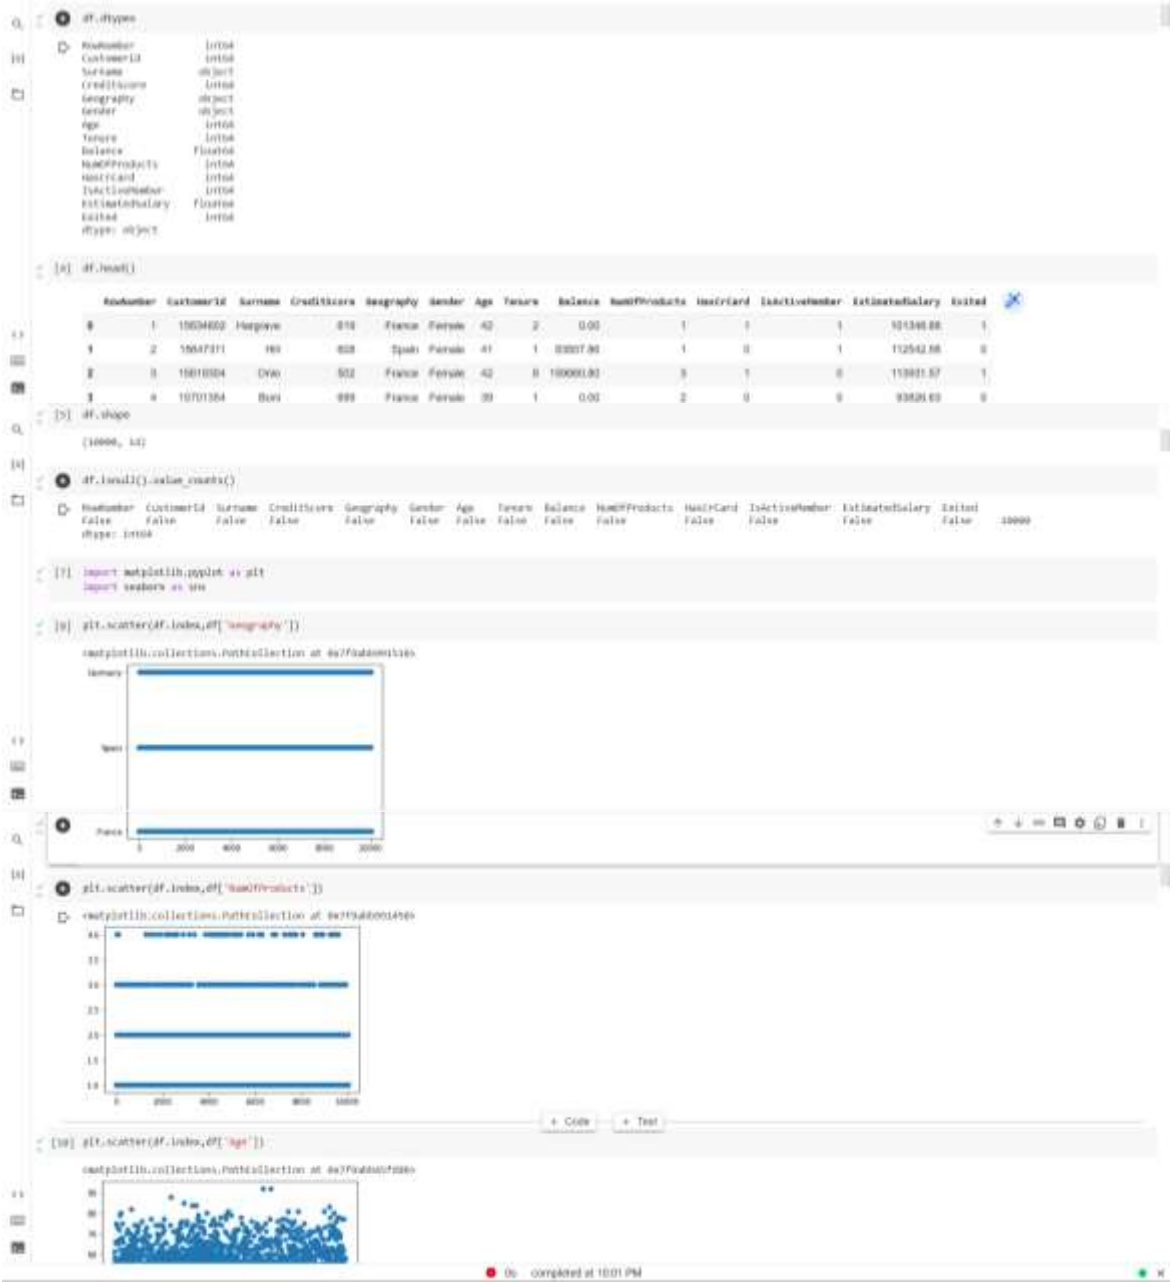
df.dtypes
df.head()
df.shape
df.isnull().sum()
df.isnull().value_counts()
import matplotlib.pyplot as plt
import seaborn as sns
plt.scatter(df.index,df['Geography'])
plt.scatter(df.index,df['NumOfProducts'])
plt.scatter(df.index,df['Age'])
+ Code
+ Text

,

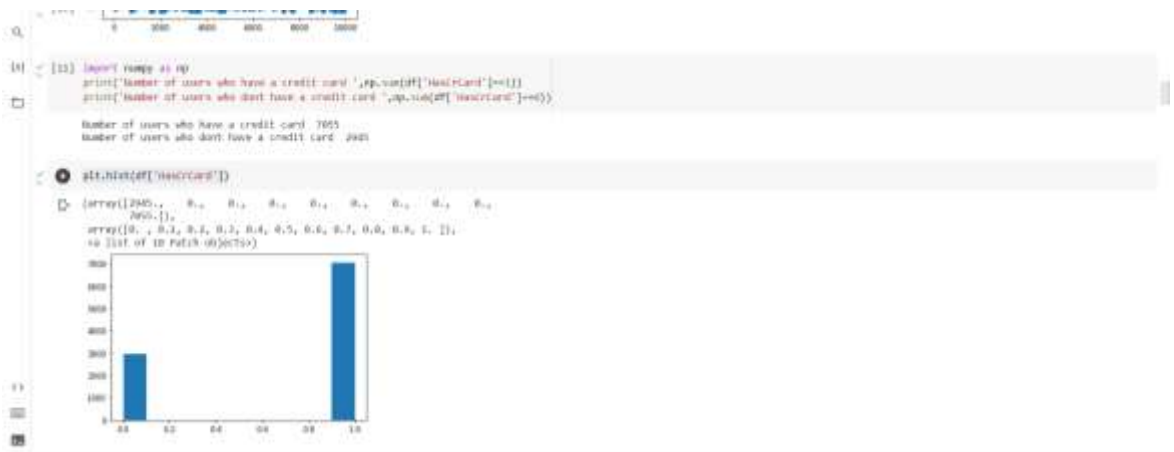
import numpy as np
print('Number of users who have a credit card ',np.sum(df['HasCrCard']==1))
print('Number of users who dont have a credit card ',np.sum(df['HasCrCard']==0))
Number of users who have a credit card 7055
Number of users who dont have a credit card 2945
plt.hist(df['HasCrCard'])

Question-3.2:

Bi-Variate

Solution:

```
fig = plt.gcf()

# Change seaborn plot size
fig.set_size_inches(20, 10)
sns.scatterplot(x=df.index,y=df['EstimatedSalary'],hue=df['HasCrCard'])

df.corr()
```

Screenshot:

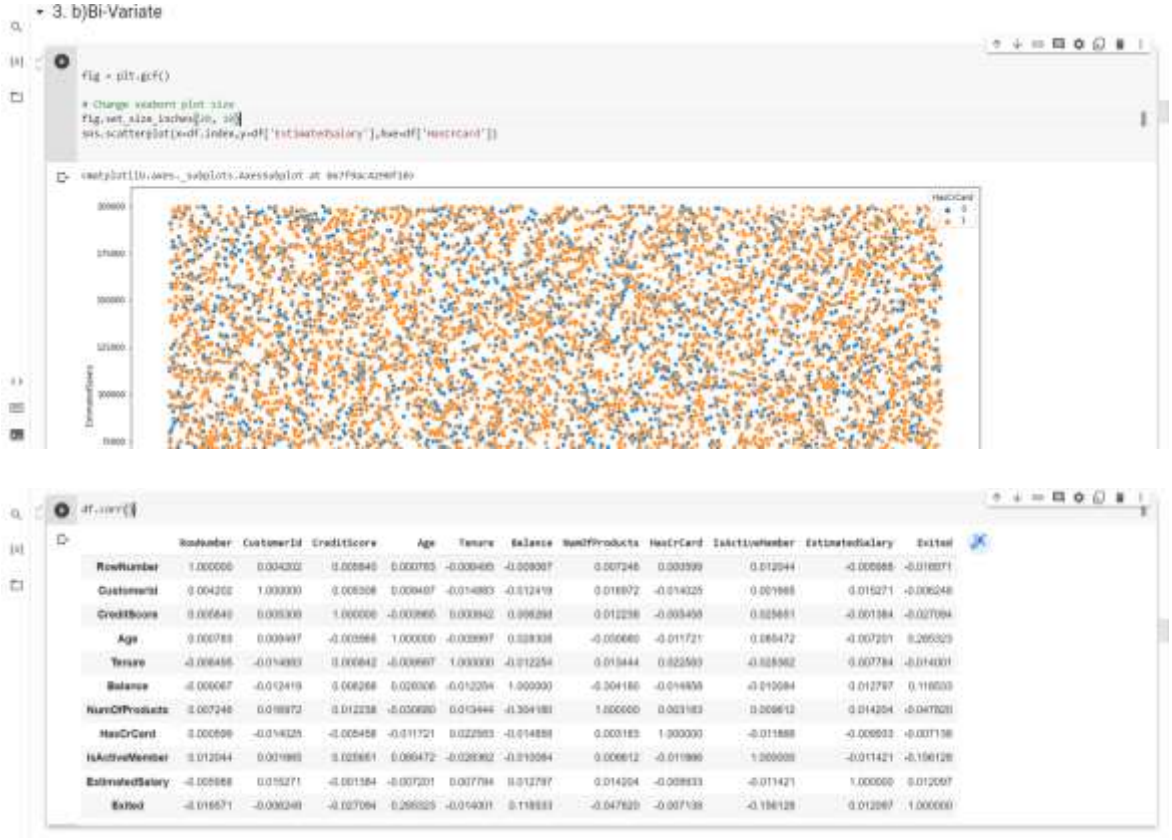

Question-3.3:

Multi Variate

Solution:

```
sns.pairplot(df)
```

Screenshot:

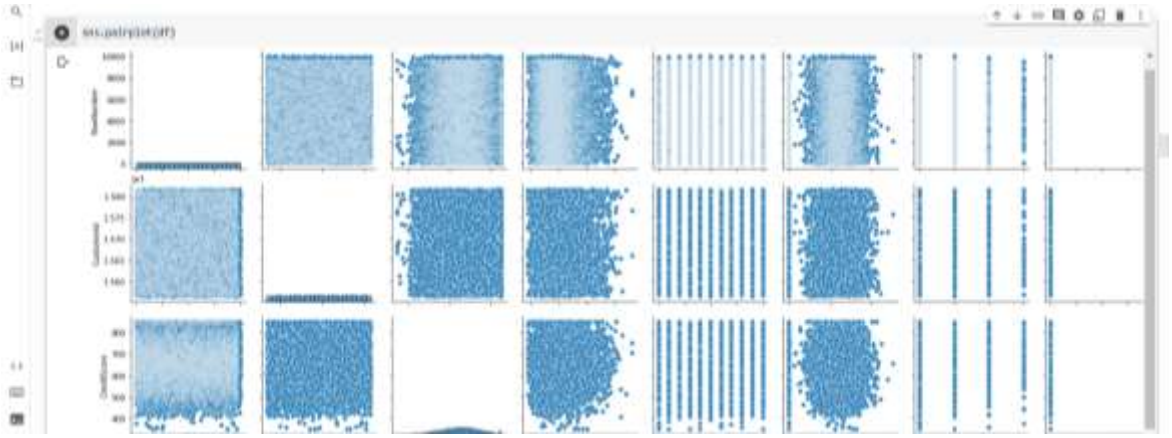

Question-4:

Perform descriptive statistics on the dataset.

Solution:

df.describe()

Screenshot:

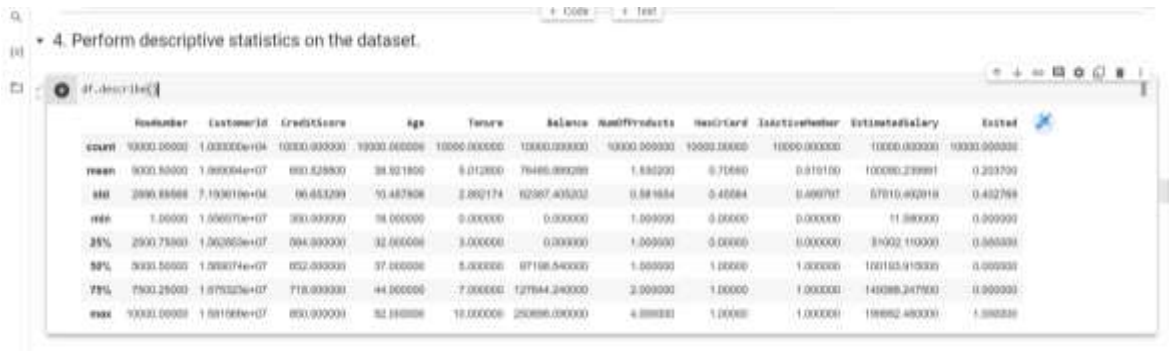

Question-5:

Handle the Missing values.

Solution:

df.isnull().value_counts()

df.notnull().value_counts()

Screenshot:

Question-6:

Find the outliers and replace the outliers

Solution:

```
df.columns

df.boxplot('HasCrCard')

numeric_col = ['Balance','CreditScore','Tenure','Age','IsActiveMember']

df.boxplot('Balance')

df.boxplot('CreditScore')
df.boxplot('Tenure')
df.boxplot('Age')
df.boxplot('IsActiveMember')
for x in ['CreditScore']:
    q75,q25 = np.percentile(df.loc[:,x],[75,25])
    intr_qr = q75-q25

    max = q75+(1.5*intr_qr)
    min = q25-(1.5*intr_qr)

    df.loc[df[x] < min,x] = np.nan
    df.loc[df[x] > max,x] = np.nan
for x in ['Age']:
    q75,q25 = np.percentile(df.loc[:,x],[75,25])
    intr_qr = q75-q25

    max = q75+(1.5*intr_qr)
    min = q25-(1.5*intr_qr)

    df.loc[df[x] < min,x] = np.nan
    df.loc[df[x] > max,x] = np.nan
df.boxplot('CreditScore')
df.boxplot('Age')
```

Screenshot:

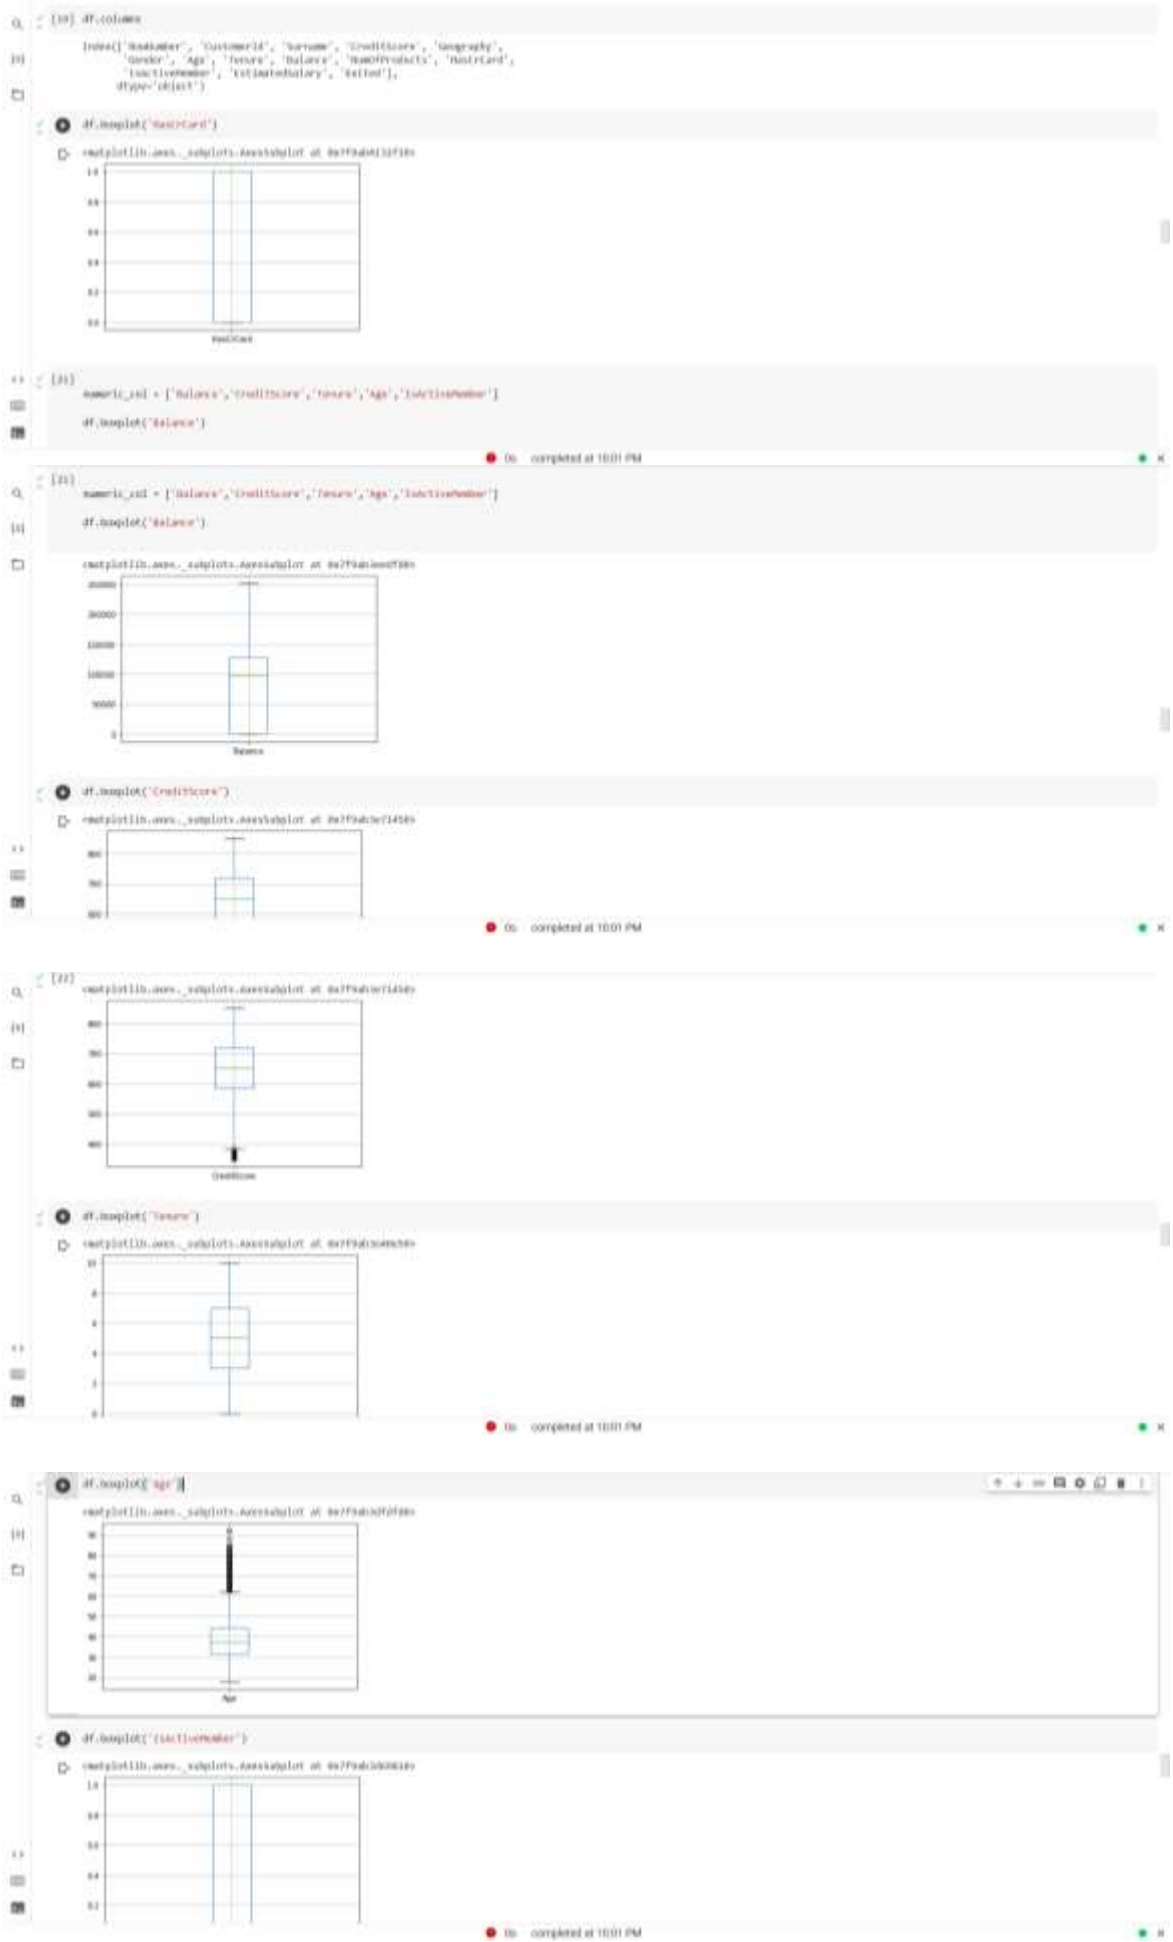

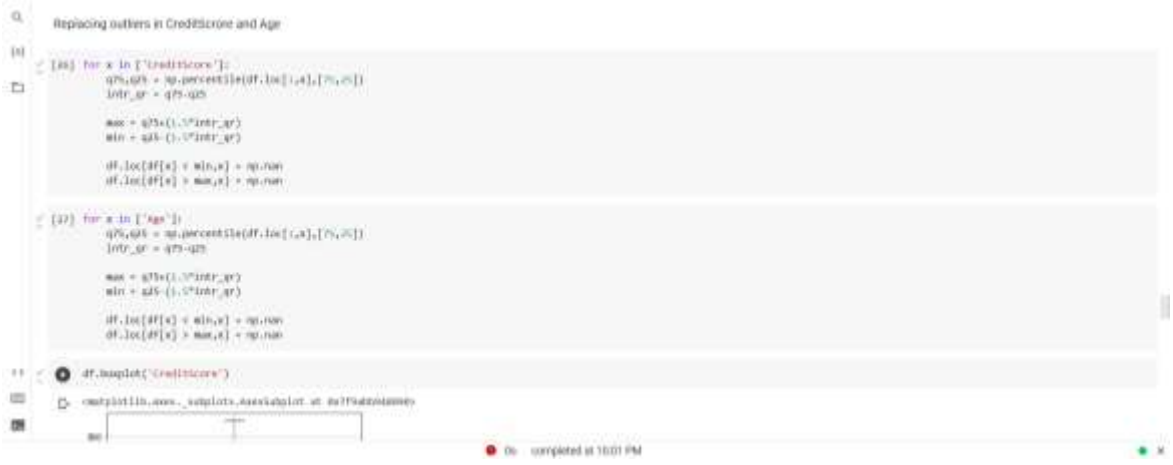

Question-7:

Check for Categorical columns and perform encoding.

Solution:

```
df.dtypes
df.head()

df_categorical=df[['Geography','Gender']]

using OneHotEncoding
from sklearn.preprocessing import OneHotEncoder
df = pd.get_dummies(df, columns = ['Geography','Gender'])
df.head()
```

Screenshot:

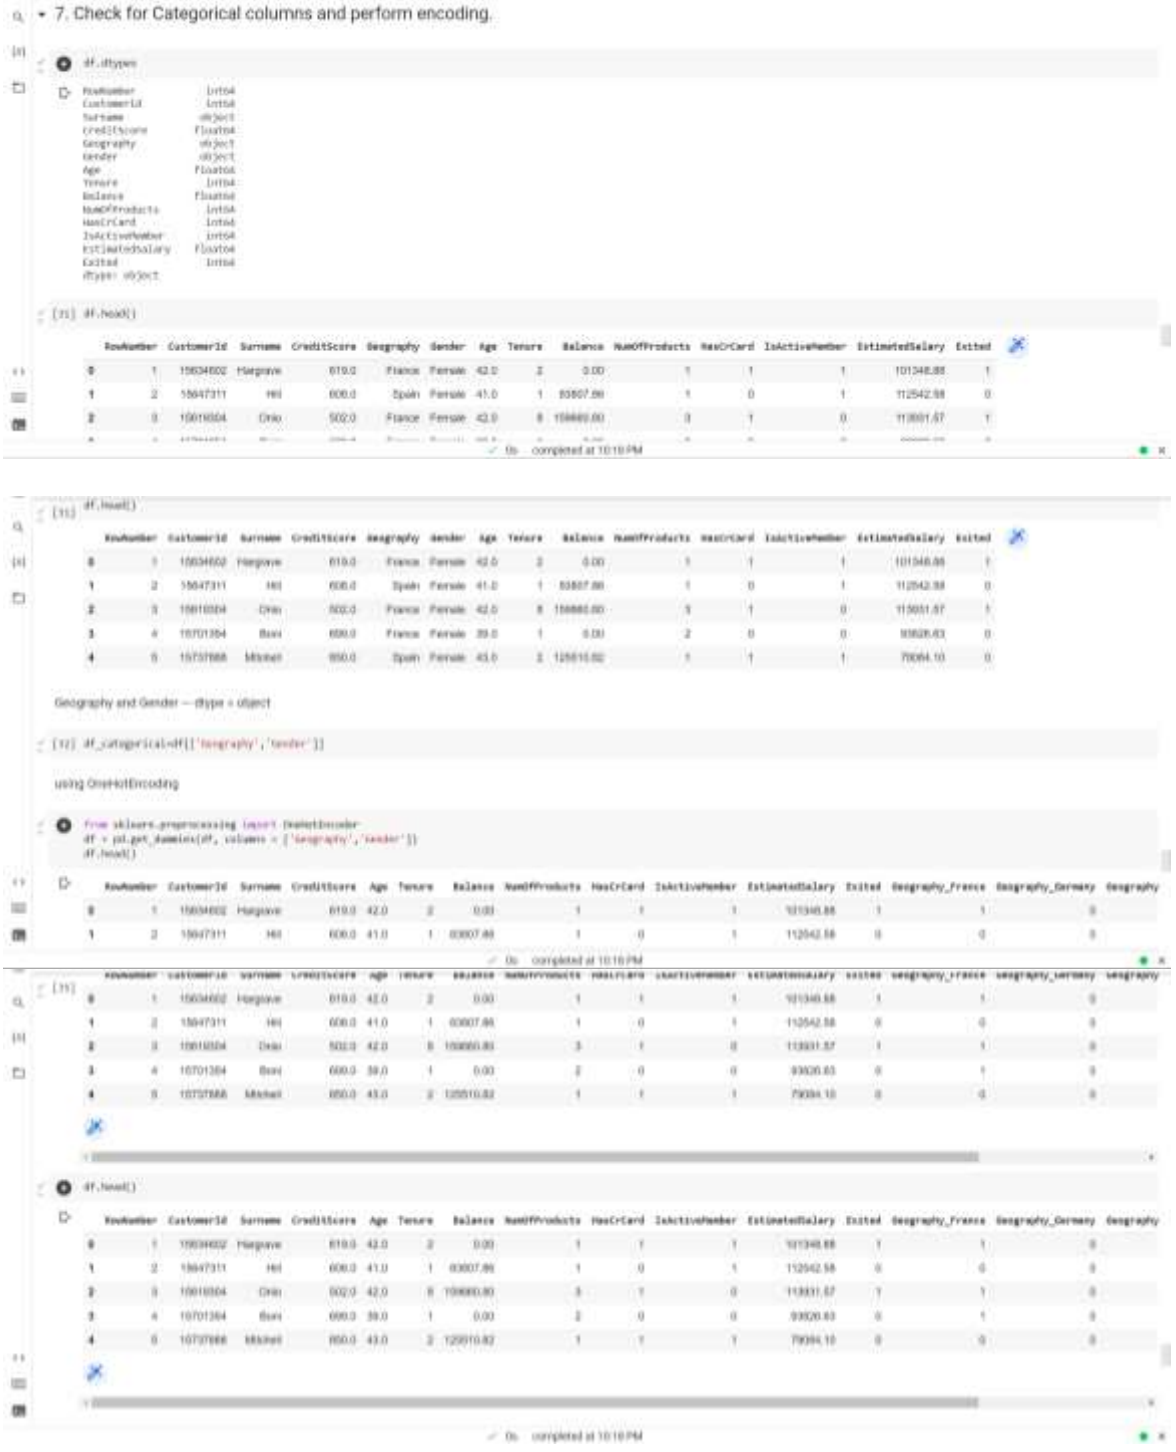

Question-8:

Split the data into dependent and independent variables

Solution:

```
df.shape
df['CustomerId'].nunique()
df['RowNumber'].nunique()
df.columns
df_independent=df[['CreditScore','Age','Tenure','Balance','NumOfProducts','IsActive
Member','EstimatedSalary','Exited', 'Geography_France', 'Geography_Germany',
      'Geography_Spain','Gender_Female', 'Gender_Male']]
df_dependent=df['HasCrCard']
```

Screenshot:

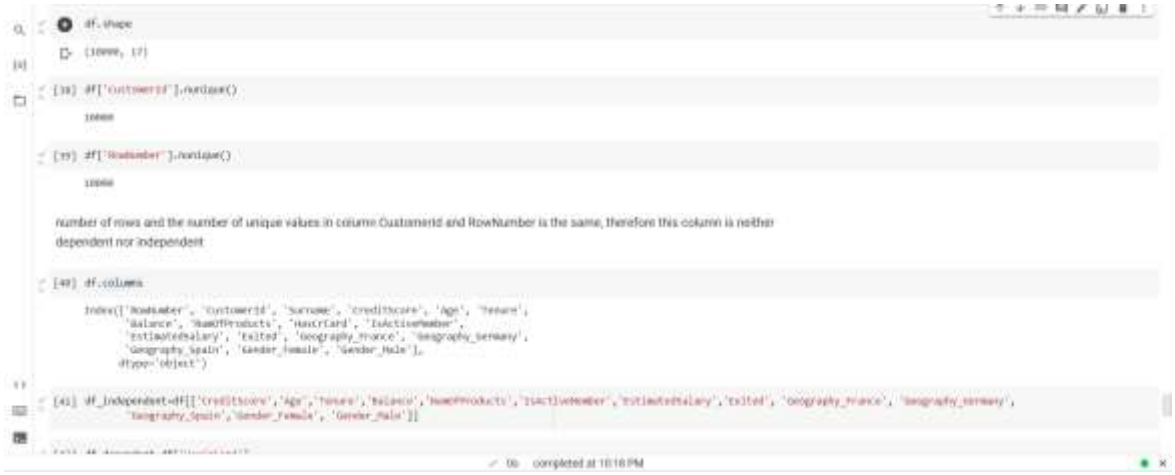

.
.
.

```
2023-09-25 00:00:00
2023-09-26 00:00:00
2023-09-27 00:00:00
2023-09-28 00:00:00
2023-09-29 00:00:00
2023-09-30 00:00:00
2023-10-01 00:00:00
2023-10-02 00:00:00
```

Question-9:

Scale the independent variables

Solution:

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

df_independent = scaler.fit_transform(df_independent)
print(df_independent)
```

Screenshot:

Question-10:

Split the data into training and testing

Solution:

```
import numpy as np
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df_independent, df_dependent, test_siz
e=0.20, random_state=42)
```

Screenshot: